

Multi-Factor ETF Strategy

Operating Document

Stuart R.
Live account: U9544585

2026-07-10

Abstract

Long-only smart-beta ETF strategy operated against a live IB Gateway account. Selects positions from a filtered universe of $\sim 2,000$ US-listed ETFs by ranking on a weighted geometric mean of four factors (momentum, quality, low-volatility, value). Top 20 held with exponential rank weights; 10% trailing stops on every position; drift-triggered rebalance ($\sim 6/\text{yr}$). This document is the operator's reading surface and takes precedence over marketing summaries; the code in `src/` is authoritative for behaviour, this document is authoritative for intent.

Contents

1	Living Document Compact	4
1.1	Purpose and Reader Contract	4
1.2	How to Keep This Document Living	4
1.3	Known Drift Between Body Sections and Live Configuration	4
1.4	Design Principles	5
2	Canonical Live Configuration	6
2.1	Portfolio Construction Parameters	6
2.2	Factor Weights (informed prior)	6
2.3	Factor Calculation Windows	7
2.4	Rebalancing	8
2.5	Risk Management	8
2.6	Optimiser Priors (used by MVO variant only)	9
3	Adversarial Review and Known Limitations	9
3.1	T1 — Material design flaws	10
3.1.1	T1.1 The value factor is fake	10
3.1.2	T1.2 Survivorship bias in the cached universe	10
3.1.3	T1.3 Look-ahead bias is presumed absent but never proved	10
3.2	T2 — Significant weaknesses	10
3.2.1	T2.1 No regime awareness	10
3.2.2	T2.2 Correlation blindness in top-N selection	11
3.2.3	T2.3 Conviction magnitude is discarded	11
3.3	T3 — Smaller nudges	11
3.3.1	T3.1 Factor lookback windows are convention, not calibrated	11
3.3.2	T3.2 Risk parameters are round numbers	11
3.3.3	T3.3 Rebalance frequency untested	11
3.4	Prioritised Fix Order	11
3.5	Strategy at a Glance	12
3.6	SPY Benchmark Comparison	12
3.7	Performance Summary	12
3.8	Capital Structure	13
4	Changes from Version 3.0 and Why They Matter	13
4.1	Universe Expansion: 1,459 → 4,953 ETFs	13
4.2	Optimizer: RankBased → Robust Mean-Variance Optimisation	14
4.3	Position Count: 20 → 30	14
4.4	Rebalancing: Bimonthly → Quarterly	15
4.5	Quantified Impact: v3.0 vs v4.0	15
5	Why the Expanded Universe Improves the Strategy	16
5.1	Better Factor Concentration	16
5.2	International Diversification	16
5.3	Lower Correlation Across Positions	16
5.4	Practical Investability	16
6	Portfolio Composition and Diversification	17
6.1	Current Holdings (as of February 2026)	17
6.2	Concentration Analysis	17
6.3	Geographic and Category Allocation	18

6.4	Rebalancing Frequency	18
7	Factor Framework	18
7.1	Factor Definitions	18
7.1.1	Momentum Factor (35%)	19
7.1.2	Quality Factor (30%)	19
7.1.3	Volatility Factor (20%)	19
7.1.4	Value Factor (15%)	19
7.2	Factor Integration	20
7.3	Justification of the Weight Vector 35/30/20/15	20
7.3.1	The empirical diagnostic and why we did not act on it	20
7.3.2	The ergodicity argument for insurance-premium factors	20
7.3.3	The Bayesian-prior critique	21
7.3.4	Why these specific integer values	21
7.3.5	When to revise	21
8	Portfolio Construction	22
8.1	Universe Filtering Pipeline	22
8.2	Robust Mean-Variance Optimisation (Default)	22
8.3	RankBased Exponential Weighting (Alternative)	23
8.4	Rebalancing Rules	23
9	Risk Management Framework	23
9.1	Automated Trailing Stops	23
9.2	Entry Stop-Loss	23
9.3	Position-Level Risk Controls	24
10	Extended Backtest Results	24
10.1	Configuration Comparison	24
10.2	Universe Scope Comparison	24
10.3	Historical Version Comparison	25
11	Implementation: Automated Pipeline	25
11.1	Architecture Overview	25
11.2	IB Gateway Integration	25
11.3	Data Collection Status	25
11.4	Trade Execution Flow	26
12	Configuration Reference	26
13	Limitations and Risks	27
14	Next Steps	27
15	Academic References	27
16	Change Log	28

1 Executive Summary

1.1 What This Is

A systematic, deterministic, long-only ETF portfolio. No LLM sets sizes or places orders. No factor timing. No sentiment trading. All numeric parameters trace to either academic literature or the calibration diagnostic under `docs/research/`.

Operated weekly: cache refresh Saturday evening, review Sunday, apply mid-week when the model proposes a rebalance and the operator approves.

1.2 Live Snapshot (2026-07-09)

Net Liquidation Value	\$390,940
Cash	\$84,552
Positions	28
Live earnings (owner's estimate)	~13–14%

Position count of 28 exceeds the target of 20 — an artefact of accumulated held positions from prior operating configurations; will normalise on the next rebalance.

1.3 Backtest Metrics (35/30/20/15 policy, top-20, RankBased exponential)

Level-1 diagnostic `docs/research/2026-07_factor_weight_diagnostic.md` on the cached ETF universe 2021-04-12 to 2026-04-16, monthly rebalance, 5 bps turnover.

Window	Sharpe	Sortino	CAGR	Max Drawdown
Full 5-year	0.47	0.60	8.07%	-9.55%
Trailing 3-year	0.96	1.22	13.32%	-9.09%

The live earnings estimate matches the 3-year backtest CAGR. The 5-year window includes the 2022 bear market drawdown; the 3-year excludes it.

1.4 Positioning

The strategy targets modest outperformance versus SPY with meaningfully lower drawdown. It is not designed to beat SPY by 5% per year. A year in which the strategy returns 12% while SPY returns 15% is a good year for the strategy *if* its drawdown is materially lower over the same period.

2 Strategy Mechanics

2.1 Universe Selection

Starting from the full IB-listed ETF universe (~5,000 tickers), the filter chain is:

1. History requirement: ≥ 252 trading days of price data.
2. Data quality: $< 10\%$ missing daily bars.
3. Leveraged/inverse exclusion: leveraged ETFs are removed by ticker prefix and issuer classification.
4. Result: ~2,000 candidates on which factor scores are computed.

2.2 Factor Framework — What the Four Factors Are and Why

The four factors below are the AQR-canonical smart-beta four: momentum, quality, low-volatility, value. They were chosen because they are the four with the largest and most robust empirical premia over the widest range of markets and time periods studied in the academic literature. AQR chose these because after decades of empirical work they are the return sources that survive out-of-sample testing, transaction costs, and factor-crowding. The strategy does not attempt to identify novel factors.

2.2.1 Momentum (target 35% weight)

$$\text{Momentum}_i = \frac{P_{i,t-21} - P_{i,t-252}}{P_{i,t-252}} \quad (1)$$

12-month return, with the most recent 21 days omitted. Winsorised at the 1st and 99th percentile.

Why: Assets that have outperformed over the prior year tend to continue outperforming over the next year. The 1-month skip avoids the short-term reversal effect that would contaminate the trend signal. Established by Jegadeesh & Titman (1993); the skip refinement from Novy-Marx (2012).

Why AQR uses it: The most-studied, most-robust factor after value, and the largest-premium factor in the last decade. Works across equity, futures, and FX. Fails during trend reversals (2022 tech drawdown, 2000 dot-com peak).

2.2.2 Quality (target 30% weight)

Composite of three z-score-normalised components on 252-day windows:

$$\text{Sharpe Ratio} = \frac{\mu_r - r_f}{\sigma_r} \sqrt{252} \quad (40\% \text{ weight}) \quad (2)$$

$$\text{Drawdown Resilience} = -1 \times \text{MaxDD} \quad (30\% \text{ weight}) \quad (3)$$

$$\text{Return Stability} = -1 \times \sigma_r \sqrt{252} \quad (30\% \text{ weight}) \quad (4)$$

Why: Assets that produce stable, high risk-adjusted returns tend to continue doing so. Established by Asness, Frazzini & Pedersen (2019).

Why AQR uses it: Quality captures the defensive-equity premium. Correlated with low-volatility but not identical — quality picks up profitability and stability of *returns*, low-vol picks up stability of *prices*. Both belong in the portfolio because they hedge different aspects of stress.

2.2.3 Low-Volatility (target 20% weight)

$$\text{Volatility}_i = \frac{1}{\sigma_{60d,i} \times \sqrt{252}} \quad (5)$$

Inverse of annualised realised volatility over 60 trading days. Higher score for lower-volatility ETFs.

Why: The low-volatility anomaly — lower-volatility assets have historically produced higher risk-adjusted returns than higher-volatility assets, contradicting the CAPM. Baker, Bradley & Wurgler (2011).

Why AQR uses it: Two reasons. First, the low-vol premium is real and persistent. Second, low-vol allocations reduce portfolio-level whipsaw, which matters more than arithmetic-mean metrics suggest in a compounding portfolio (see §??).

2.2.4 Value (target 15% weight) — currently degraded

$$\text{Value}_i = -1 \times \text{Expense Ratio}_i \quad (6)$$

This is not a value factor. It is a low-fee preference. A real value factor uses fundamental multiples (P/E, P/B, dividend yield, EBITDA/EV). The current implementation was chosen for simplicity and is on the fix list at §?? T1.1. Until fixed, the 15% weight allocated to value is buying almost no cross-sectional signal.

Why AQR uses value: It is the oldest and most-studied factor (Fama & French 1993). Tends to work when growth and momentum break (2000–2002, 2022). Its absence here is the strategy’s largest current gap.

2.3 Factor Integration

$$\text{Score}_i = \text{Mom}_i^{0.35} \times \text{Qual}_i^{0.30} \times \text{Vol}_i^{0.20} \times \text{Val}_i^{0.15} \quad (7)$$

Each factor is first converted to a percentile rank in $[0,1]$. The integrated score is the weighted geometric mean of ranks.

Why geometric mean, not arithmetic: Geometric-mean integration penalises weakness on any one factor. An ETF in the 90th percentile on three factors but the 20th on one ranks below an ETF in the 70th percentile on all four. This “multi-factor consistency” requirement — explicitly the key insight from AQR’s factor-integration research — rejects one-trick candidates in favour of broadly-strong candidates. Arithmetic integration would allow one strong factor to compensate for weakness on the others.

2.4 Portfolio Construction

1. Score every ETF in the filtered universe.
2. Rank descending.
3. Take the top 20.
4. Assign exponential rank weights: $w_i \propto e^{-\text{rank}_i/20}$, normalised to sum to 1.
5. Apply weight bounds: minimum 2%, maximum 15%.

Exponential weighting expresses conviction — the top-ranked position gets meaningfully more weight than the twentieth — while the bounds prevent single-position concentration. The alternative optimiser (`OPTIMIZER.TYPE = "mvo"`) is available in the code but not the live default.

2.5 Rebalancing

1. Scheduled: bimonthly (~ 6 times per year).
2. Drift-triggered: if any position’s actual weight deviates from target by $>5\%$, rebalance.
3. In practice, drift triggers dominate; scheduled rebalances rarely fire independently.

Drift threshold is deliberately wide. Tighter thresholds cause more trading on smaller deviations — textbook whipsaw. See §??.

2.6 Risk Management

Every entry BUY places a 10% trailing stop, GTC, on the position. Trailing stops ratchet up as the position appreciates. On any 10% pullback from the running high, the stop triggers.

Entry stop-loss of 12% (a bit wider than the trailing stop) accommodates initial-position volatility before the trailing high-water mark establishes.

Cash reserve of \$70,000 is held unallocated. This is a static allocation, not a regime-responsive hedge.

2.7 Explicit Assumptions

The strategy assumes:

- Factor premia (momentum, quality, low-vol, value) will continue to exist over the holding horizon. They may attenuate but do not disappear.
- The rank structure of the universe changes slowly enough that monthly rebalancing captures the signal without significant slippage.
- Retail transaction costs at IB (~ 5 bps effective) are small enough not to erode the premium.
- Long-only holds; no short exposure to hedge market drawdown.
- No leverage.
- Trailing stops execute at or near their trigger prices. Gap-down scenarios (weekend news, halt reopens) are unhedged.
- Position selection ignores tax lot-level implications. Tax-lot management is out of scope.

2.8 Design Principles

Three non-negotiable filters govern every future proposal:

1. **Deterministic decides.** Every number reaching the broker comes from code in `src/`. No LLM output influences a size or weight.
2. **No magical thinking on knobs.** Every parameter traces to research or a repo-recorded backtest. Knobs whose provenance is intuition go on the research backlog.
3. **Slowly varying only.** Whipsaw is evidence of overfitting, not alpha. Long lookbacks are preferred over short. Drift thresholds err high. Regime signals (when added) require hysteresis and minimum dwell time. Factor timing is rejected on principle.

3 Parameters and Weights

3.1 The Factor Weights

The weight vector is momentum = 35%, quality = 30%, low-volatility = 20%, value = 15%. This is an *informed prior*, deliberately regularised away from any in-sample optimum.

3.1.1 Why not the in-sample winner?

The July 2026 diagnostic (grid search over the three-factor simplex, momentum/quality/low-vol, on the cached 2021–2026 history) produced in-sample winners uniformly momentum-heavy: 80–90% momentum, 0–30% quality, 10–20% low-vol. Sharpe up to 1.30 on the 3-year window versus ≈ 0.96 for the current policy.

This finding was deliberately not acted on. Three reasons:

1. **Single-regime evidence.** 2021–2026 was overwhelmingly momentum-friendly (post-COVID rally, mega-cap tech, AI narrative). Both the 5-year and 3-year windows overlap and sample the same regime. This is not an independent invariance test.
2. **Non-ergodicity.** Equity returns are multiplicative. For a single investor on a single path, the arithmetic mean of returns diverges from the time-average growth rate they actually experience (Peters 2019). A 50% drawdown requires 100% subsequent gain to recover, not 50%. In-sample-Sharpe optimisation understates this asymmetry. Factors that reduce drawdown variance (quality, low-vol) buy protection against the multiplicative losses a single-path investor most needs to avoid. These are insurance premia, not underperforming factors.

3. **Confirmation bias.** The operator is explicitly on record as most concerned about an AI/tech valuation regime unwind. The mechanical winner (90% momentum) is exactly the exposure that would be worst affected in that scenario. Adopting the in-sample winner would be confirming a bias against the operator’s stated risk view.

3.1.2 The Constraint Framework

The 35/30/20/15 vector satisfies three constraints:

1. **Ordering: momentum > quality > low-vol > value.** Reflects the ordering of the empirical premia in the AQR literature and the compounding effect that the value implementation is currently a low-fee proxy rather than a true value factor.
2. **No factor below 10%.** A factor below 10% receives such a small weight in the weighted geometric mean that it fails to discriminate between candidates. The multi-factor-consistency requirement collapses to a three-factor screen.
3. **No factor above 40%.** A factor above 40% dominates the integrated score to the point that single-factor selection would produce nearly the same portfolio. This defeats the multi-factor rationale and re-introduces single-factor tail risk.

Within those constraints the exact values are round numbers chosen for auditability. They can be reproduced in one sentence to a future maintainer: “momentum gets the largest weight but capped, quality gets the second-largest, low-vol gets a solid third to hedge whipsaw, value gets the smallest until we fix the proxy.”

3.1.3 When To Revise

Revise the weights when at least one of the following is available:

- Backfilled history to pre-2020 that permits a non-overlapping regime comparison (2011–2015 vs 2021–2026).
- A working value factor using fundamentals (P/E, P/B, dividend yield) rather than expense ratio.
- A deliberate operator-declared factor-level sentiment override, committed to `configs/etf_smart_beta.toml` with a dated research note.

3.2 Portfolio Construction Parameters

Parameter	Value	Justification
<code>num_positions</code>	20	Target holding count. Prior operating configurations without a fixed target produced degenerate 1–2 position portfolios when factor scores were bimodal; the fixed target of 20 forces diversification.
<code>min_weight</code>	2%	Position floor. Ensures a full-basket portfolio uses at most 40% of NAV at floor, leaving room for overweight tilt on high-conviction positions.
<code>max_weight</code>	15%	Position cap. Prevents single-position concentration under the RankBased scheme; the top-ranked ETF cannot receive more than 15% of NAV.
<code>tax_status</code>	<code>taxable</code>	Signals real capital-gains consequences to any future tax-lot / wash-sale logic.

3.3 Factor Calculation Windows

Parameter	Value	Justification
Momentum lookback	252 days	Jegadeesh & Titman (1993). 12-month formation window.
Momentum skip	21 days	Novy-Marx (2012). Avoid short-term reversal.
Quality lookback	252 days	AQR convention.
Value lookback	252 days	Unused while value is expense-ratio proxy.
Volatility lookback	60 days	~3 months of daily returns. Standard practice.

3.4 Rebalancing Parameters

Parameter	Value	Justification
Scheduled frequency	bimonthly (~6/yr)	Balances signal freshness against turnover.
Drift threshold	5%	Round number, defensible per Donohue-Yip (2003). Widening (7.5%, 10%) considered before tightening.

3.5 Risk Management Parameters

Parameter	Value	Justification
Entry stop-loss	12%	Slightly wider than trailing stop to accommodate initial position volatility.
Trailing stop	10%	Sized to sit above ordinary noise ($\approx 2-3\sigma$ daily) so it triggers only on genuine trend breaks.
Cash reserve	\$70,000	$\approx 18\%$ of current NAV. Static, not regime-responsive. Fix on the list (§?? T2.1).

3.6 Optimiser Priors (MVO variant only)

The live default optimiser is RankBased. The values below apply only when `OPTIMIZER_TYPE = "mvo"`.

Parameter	Value	Justification
Risk aversion λ	1.5	Standard MVO. Meucci retail range 2–4; 1.5 is risk-tolerant end.
Robustness penalty γ	0.7	Axioma-style, aggressive. Chosen default.
Turnover penalty	0.2	Medium. Chosen default.

4 Current Flaws

Adversarial review, tiered by impact. Items remain here until closed; closed items move to the Change Log with the closing commit.

4.1 Tier 1 — material design flaws

- **T1.1 The value factor is fake.** `SimplifiedValueFactor` returns $-1 \times$ expense ratio. This is a fee preference, not a value signal. The 15% weight buys almost nothing. The natural regime hedge that value provides (2000–2002, 2022) is absent from the model. **Highest-impact single fix.** Replace with fundamentals composite (P/E, P/B, dividend yield); IB Fundamentals API preferred over yfinance.
- **T1.2 Survivorship bias in the cached universe.** The per-ticker parquet cache contains only ETFs that still trade. Delisted / liquidated ETFs are absent. Every backtest overstates historical performance by the average return of survivors. Mitigation: 1–2% CAGR haircut on all backtested figures.
- **T1.3 Look-ahead bias presumed absent, never proved.** No pinning test asserts that factor calculations at time t use only data with index $< t$. If any factor uses a window that includes $t+1$, the backtest is silently inflated. Fix: add a shift-and-recompute test. Two hours' work.

4.2 Tier 2 — significant weaknesses

- **T2.1 No regime awareness.** Strategy runs identically in bull, bear, and crisis regimes. Long-only, concentrated, static 22% cash reserve, 10% trailing stops. No signal to de-risk in rising-VIX / breadth-collapsing environments. Fix: add SPY-vs-200d moving average and VIX overlay; when both signals point to stress, shift 100% target equity to 60%. Must include hysteresis and ≥ 30 -day dwell time in either regime (per Principle 3).
- **T2.2 Correlation blindness in top-N selection.** The integrated score picks the top 20 ETFs by factor score without checking their pairwise correlations. Current live holdings read as a broad international dividend / value / quality tilt with likely 0.6–0.8 pairwise correlations. Effective diversification is materially less than 20 positions suggests. Fix: take top 40 by score, cluster on 6-month correlations, pick top-scored representative per cluster until 20 positions filled.
- **T2.3 Conviction magnitude discarded.** The integrated score's magnitude (0.95 vs 0.62) is not used in the optimiser; only rank order matters. Exponential rank weighting mitigates partially. Fix: pass score magnitude as expected-return alpha to the optimiser.

4.3 Tier 3 — smaller nudges

- **T3.1 Factor lookback windows are convention.** 252 / 21 / 60 chosen by AQR convention, never tested for optimality in this universe.
- **T3.2 Risk parameters are round numbers.** 10%, 12%, \$70k, 5% drift, 2–15% weight bounds — all defensible, all uncalibrated. Sensitivity grid would locate them on the risk-return frontier.
- **T3.3 Rebalance frequency untested.** Bimonthly vs quarterly turnover trade-off never explicitly measured on this universe.

4.4 Priority Order

1. T1.3 (look-ahead pinning test) — 2 hours. Proves the model is not silently a lie.
2. T1.1 (real value factor) — 1–2 evenings. Single highest-impact fix.
3. T2.1 (regime overlay) — 1 evening. Drawdown protection before live orders resume.
4. T2.2 (correlation clustering) — 1 evening. Effective diversification.
5. T1.2 (survivorship haircut) — 1 hour. Honesty in reporting.
6. T2.3, T3. — as time permits.

5 Change Log

Chronological, most recent first. Each entry cites the closing commit SHA(s). Every commit touching `configs/etf_smart.beta.toml`, `src/factors/`, `src/portfolio/`, or the execution scripts must append an entry.

2026-07-10 — Full rewrite

Document rewritten from the ground up. Historical version-comparison content (v3 → v4 migration story, v4.2 backtest tables, Feb-2026 holdings snapshot) removed. New structure: Executive Summary → Strategy Mechanics → Parameters and Weights → Current Flaws → Change Log. Factor definitions and integration rationale promoted into the Mechanics section; the constraint framework (10%–40% weight bounds, factor ordering) promoted into the Parameters section. Length approximately halved.

2026-07-09 — Adversarial review + weight justification (superseded by rewrite)

Adversarial review commissioned and recorded. Ergodicity-economics justification for retaining the 35/30/20/15 weight vector recorded in the Factor Framework section. Design Principles (deterministic decides, no magical thinking, slowly varying) recorded as non-negotiable filters. Level-1 factor-weight diagnostic produced at `docs/research/2026-07_factor_weight_diagnostic.md`. Commits d9826dc, 2198273, 6a14212.

2026-02-15 — Prior v4.2 design (historical, retained for context)

Prior design documented 30 positions, 3–8% weight bounds, quarterly rebalance, Robust MVO default. Backtested 12.1% CAGR / 0.64 Sharpe over 2021–2026. Live configuration diverged from this design in three places (position count, weight bounds, default optimiser) and the 2026-07 rewrite reconciled the document to what is actually running.